

LP1 Assignment HPC H1

Parallel Computing using CUDA

Date - 16th August, 2020.

Assignment Number - HPC H1

Title

Parallel Computing using CUDA

Problem Definition

1. Implement Parallel Reduction using Min, Max, Sum and Average operations.
2. Write a CUDA program that, given an N-element vector, find-
 - a. The maximum element in the vector
 - b. The minimum element in the vector
 - c. The arithmetic mean of the vector
 - d. The standard deviation of the values in the vector

Test for input N and generate a randomized vector V of length N (N should be large). The program should generate output as the two computed maximum values as well as the time taken to find each value.

Learning Objectives

- Learn parallel decomposition of problems.
- Learn parallel computing using CUDA

Learning Outcomes

I will be able to decompose problems into subproblems, to learn how to use GPUs, to learn to solve sub problems using threads on GPU cores.

Software Packages and Hardware Apparatus Used

- Operating System : 64-bit Ubuntu 18.04
- Browser : Google Chrome
- Programming Language : C++, Python 3
- Jupyter Notebook Environment : Google Colaboratory

Related Mathematics

Mathematical Model

Let S be the system set:

$$S = \{s; e; X; Y; F_{me}; F_f; DD; NDD; F_c; S_c\}$$

s=start state

e=end state

X=set of inputs

$$X = \{X_1\}$$

where X_1 = Vector

Y= Output Set

$$Y = \{Y_1, Y_2, Y_3, Y_4, Y_5, Y_6\} \text{ where}$$

Y_1 = Minimum element

Y_2 = Maximum element

Y_3 = Average of all elements

Y_4 = Sum of all elements

Y_5 = Variance

Y_6 = Standard Deviation

F_{me} is the set of main functions

$$F_{me} = \{f_0\} \text{ where}$$

F_0 = output display function

Ff is the set of friend functions

Ff = {f1,f2,f3,f4,f5,f6,f7,f8,f9} where

f1 = function to find Minimum

f2 = function to find Maximum

f3 = function to find Sum

f4 = function to find Mean

f5 = function to find Variance

f6 = function to find Standard Deviation

f7 = function to write CUDA code from Google Colab Cell
into a file and compile the code

f8 = function to initialize vector with random elements

f9 = function to display vector

DD = Deterministic Data

Input Vector X1

NDD=Non-deterministic data

No non deterministic data

Fc =failure case:

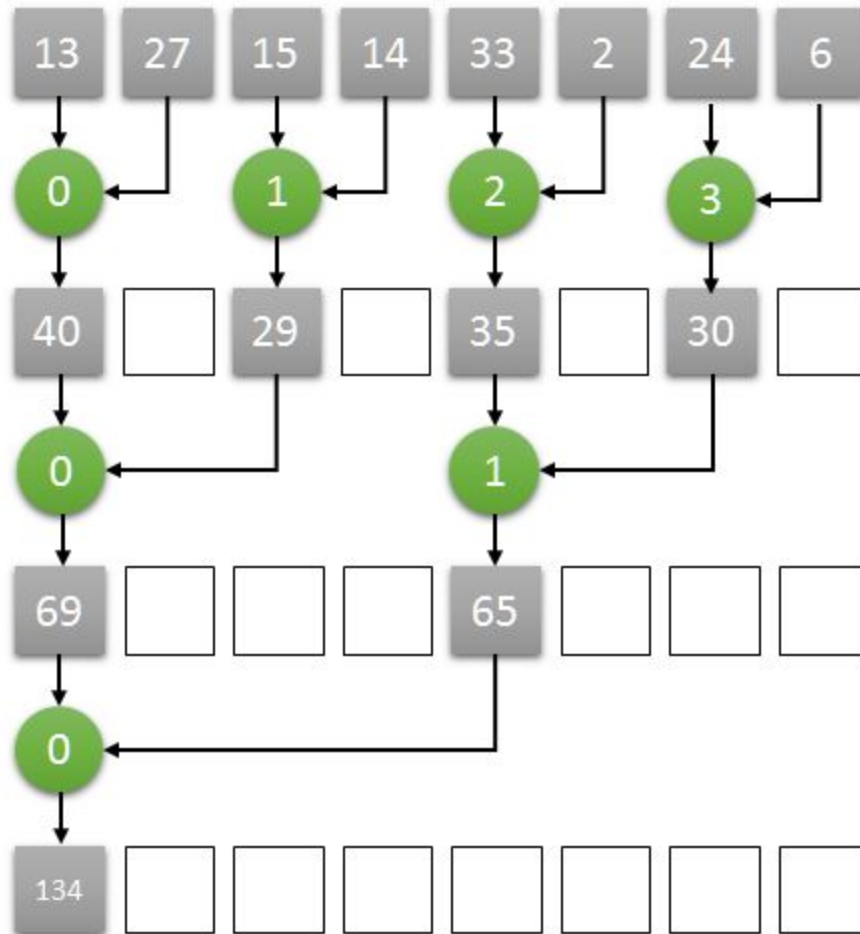
No failure case identified for this application

Concepts related Theory

Parallel computing is a type of computing architecture in which several processors or processes execute or process an application or computation simultaneously. The primary objective of parallel computing is to increase the available computation power for faster application processing or task resolution

Diagram for Parallel Reduction

Sum Function Demonstrated



Algorithms

Algorithm to get sum of all the elements

1. Let N be the number of elements
2. threadCount is N/2 which compares two elements at a time
3. threadCount := N/2
4. stepSize := 1
5. Each thread stores the sum of the two elements in place of first element
6. For each iteration
 - a. Select two consecutive elements
 - b. Compute their sum
 - c. Store the sum in place of first array
 - d. threadCount := threadCount/2
 - e. stepSize := stepSize*2

7. return first element of the array

Algorithm to get maximum element

1. Let N be the number of elements
2. threadCount is N/2 which compares two elements at a time
3. threadCount := N/2
4. stepSize := 1
5. Each thread stores the sum of the two elements in place of first element
6. For each iteration
 - a. Select two consecutive elements
 - b. Compute the maximum between the two
 - c. Store the sum in place of first array
 - d. threadCount := threadCount/2
 - e. stepSize := stepSize*2
7. return first element of the array

Algorithm to get minimum element

1. Let N be the number of elements
2. threadCount is N/2 which compares two elements at a time
3. threadCount := N/2
4. stepSize := 1
5. Each thread stores the sum of the two elements in place of first element
6. For each iteration
 - a. Select two consecutive elements
 - b. Compute the minimum between the two
 - c. Store the sum in place of first array
 - d. threadCount := threadCount/2
 - e. stepSize := stepSize*2
7. return first element of the array

Algorithm to get variance

1. Let N be the number of elements
2. threadCount is N/2 which compares two elements at a time
3. threadCount := N/2
4. stepSize := 1
5. Each thread stores the sum of the two elements in place of first element
6. Mean Square all elements in place
7. Get sum of all elements, store it in var

8. $\text{Var} := \text{var}/N$
9. return

Greek Symbols in Probability and Statistics

- 1) μ represents the population [mean](#)
- 2) σ represents the population [standard deviation](#)

Source Code

```
#include<iostream>
#include<math.h>

#define n 8

using namespace std;

//Kernel Functions

__global__ void minimum(int *input) {

    int threadId = threadIdx.x;
    int stepSize = 1;
    int threadCount = blockDim.x;

    while(threadCount>0) {
        if(threadId < threadCount) {
            int first = threadId*stepSize*2;
            int second = first + stepSize;
            //Modify Array In Place
            if(input[second] < input[first])
                input[first] = input[second];
        }
        stepSize *=2;
        threadCount /= 2;
    }
}
```

```

__global__ void maximum(int *input) {
    int threadId = threadIdx.x;
    int stepSize = 1;
    int threadCount = blockDim.x;

    //Array is updated Inplace
    while(threadCount>0) {
        if(threadId < threadCount) {
            int first = threadId*stepSize*2;
            int second = first + stepSize;
            //Modify Array In Place
            if(input[second] > input[first])
                input[first] = input[second];
        }
        stepSize <=< 1;
        threadCount >>= 1;
    }
}

__global__ void sum(int *input) {
    const int threadId = threadIdx.x;
    int stepSize = 1;
    int threadCount = blockDim.x;

    while(threadCount > 0) {
        if(threadId < threadCount) {
            int first = threadId * stepSize * 2;
            int second = first + stepSize;
            //Modify Array In Place
            input[first] += input[second];
        }
        stepSize <=< 1;
        threadCount >>= 1;
    }
}

```

```

__global__ void meanSquared(float *input, float mean) {
    input[threadIdx.x] -= mean;
    input[threadIdx.x] *= input[threadIdx.x];
}

```

```

__global__ void sum(float *input) {
    int threadId = threadIdx.x;
    int stepSize = 1;
    int threadCount = blockDim.x;

    while(threadCount > 0) {
        if(threadId < threadCount) {
            int first = threadId * stepSize * 2;
            int second = first + stepSize;
            //Modify array in place
            input[first] += input[second];
        }
        stepSize <= 1;
        threadCount >>= 1;
    }
}

```

//Host Functions

```

void initialize_vector(int *input, int size) {
    for(int i=0; i<size; i++) {
        input[i] = rand()%500;
    }
}

```

```

void display_vector(int *input, int size) {
    if(size==0){
        cout<<"[]";
        return;
    }
}

```



```

    cout<<"["<<input[0];
    for(int i=1; i<size; i++) {
        cout<<" , "<<input[i];
    }
    cout<<" ]";
}

int getMinimum(int* arr_d, int* arr, int size){
    int result;
    cudaMemcpy(arr_d, arr, size, cudaMemcpyHostToDevice);
    minimum<<<1,n/2>>>(arr_d);
    //Copying Just the first Element of the array
    cudaMemcpy(&result, arr_d, sizeof(int), cudaMemcpyDeviceToHost);
    return result;
}

int getMaximum(int* arr_d, int* arr, int size){
    int result;
    cudaMemcpy(arr_d, arr, size, cudaMemcpyHostToDevice);
    maximum<<<1,n/2>>>(arr_d);
    //Copying Just the first Element of the array
    cudaMemcpy(&result, arr_d, sizeof(int), cudaMemcpyDeviceToHost);
    return result;
}

int getSum(int* arr_d, int* arr, int size){
    int result;
    cudaMemcpy(arr_d, arr, size, cudaMemcpyHostToDevice);
    sum<<<1,n/2>>>(arr_d);
    //Copying Just the first Element of the array
    cudaMemcpy(&result, arr_d, sizeof(int), cudaMemcpyDeviceToHost);
    return result;
}

float getAverage(int* arr_d, int* arr, int size){
    float result = float(getSum(arr_d,arr,size))/n;
    return result;
}

```

```

}

float getVariance(int* arr_d, int* arr, int size){

    //int n = size/sizeof(int);

    float mean = getAverage(arr_d, arr, size);

    float *arr_float;
    float *arr_std, result;

    arr_float = (float *)malloc(n*sizeof(float));

    cudaMalloc((void **)&arr_std, n*sizeof(float));

    for(int i=0; i<n; i++)
        arr_float[i] = float(arr[i]);

    //Mean Squared
    cudaMemcpy(arr_std, arr_float, n*sizeof(float),
cudaMemcpyHostToDevice);
    meanSquared <<<1,n>>>(arr_std, mean);
    cudaMemcpy(arr_std, arr_float, n*sizeof(float),
cudaMemcpyDeviceToHost);

    //Add Mean Squared of all the elements
    sum<<<1,n/2>>>(arr_std);
    cudaMemcpy(&result, arr_std, sizeof(float), cudaMemcpyDeviceToHost);

    result /= n;

    cudaFree(arr_std);

    return result;

}

```

```
float getStdDeviation(int* arr_d, int* arr, int size){  
    float result = sqrt(getVariance(arr_d,arr,size));  
    return result;  
}
```

```
int main() {  
  
    //Host Variable  
    int *arr;  
  
    //Device Variable  
    int *arr_d;  
  
    int size = n*sizeof(int);  
  
    //Allocate Memory to Host Variable  
    arr = (int *)malloc(size);  
  
    initialize_vector(arr, n);  
    cout<<endl<<"Vector - ";  
    display_vector(arr, n);  
  
    //Allocate Memory to Device Variable  
    cudaMalloc((void **)&arr_d, size);  
  
    //Output Variables  
    int min, max, sum;  
    float avg, var, stddev;  
  
    //Host Function Calls - They Launch the Kernel  
    min = getMinimum(arr_d, arr, size);  
    max = getMaximum(arr_d, arr, size);  
    sum = getSum(arr_d, arr, size);  
    avg = getAverage(arr_d, arr, size);  
    var = getVariance(arr_d, arr, size);  
    stddev = getStdDeviation(arr_d, arr, size);  
}
```

```
//Output to the console
cout<<endl<<"\nMinimum - "<<min;
cout<<endl<<"Maximum - "<<max;
cout<<endl<<"Sum - "<<sum;
cout<<endl<<"Average - "<<avg;
cout<<endl<<"Variance - "<<var;
cout<<endl<<"Standard Deviation - "<<stddev;

//Freeing space
free(arr);
cudaFree(arr_d);

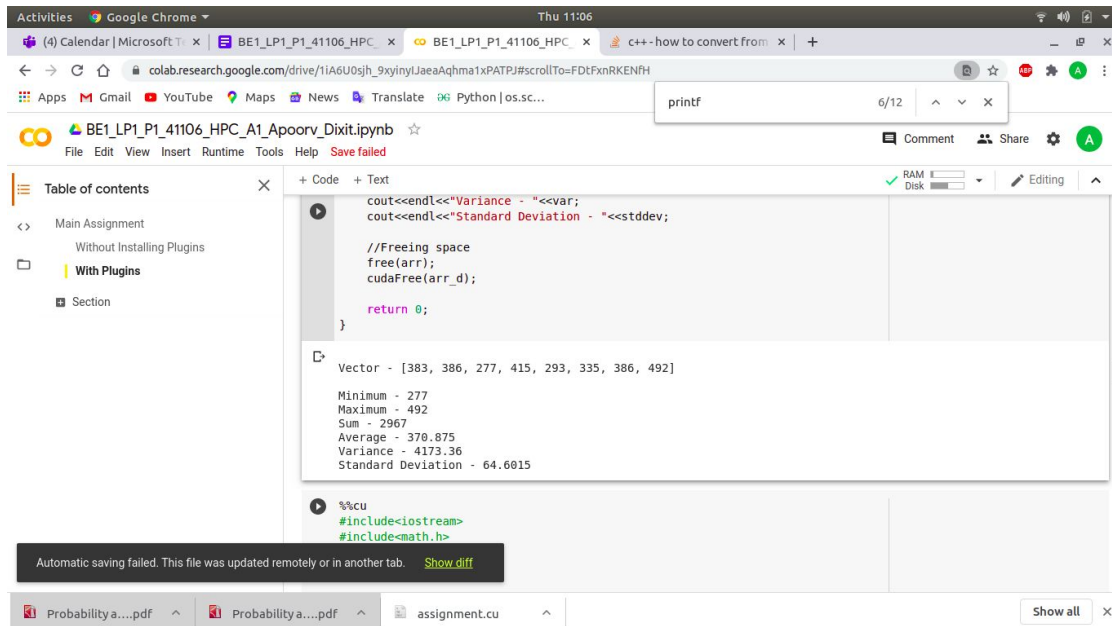
return 0;
}
```

Google Colab Notebook Link

https://colab.research.google.com/drive/1iA6U0sjh_9xyinyIJaeaAqhma1xPATPJ?usp=sharing

Output Screenshots

Min, Max, Sum, Average, Variance, Standard Deviation



The screenshot shows a Jupyter Notebook interface with a C++ code cell. The code calculates the minimum, maximum, sum, average, variance, and standard deviation of an array. The output is displayed as a text box below the code cell.

```
cout<<endl<<"Variance - "<<var;
cout<<endl<<"Standard Deviation - "<<stddev;

//Freeing space
free(arr);
cudaFree(arr_d);

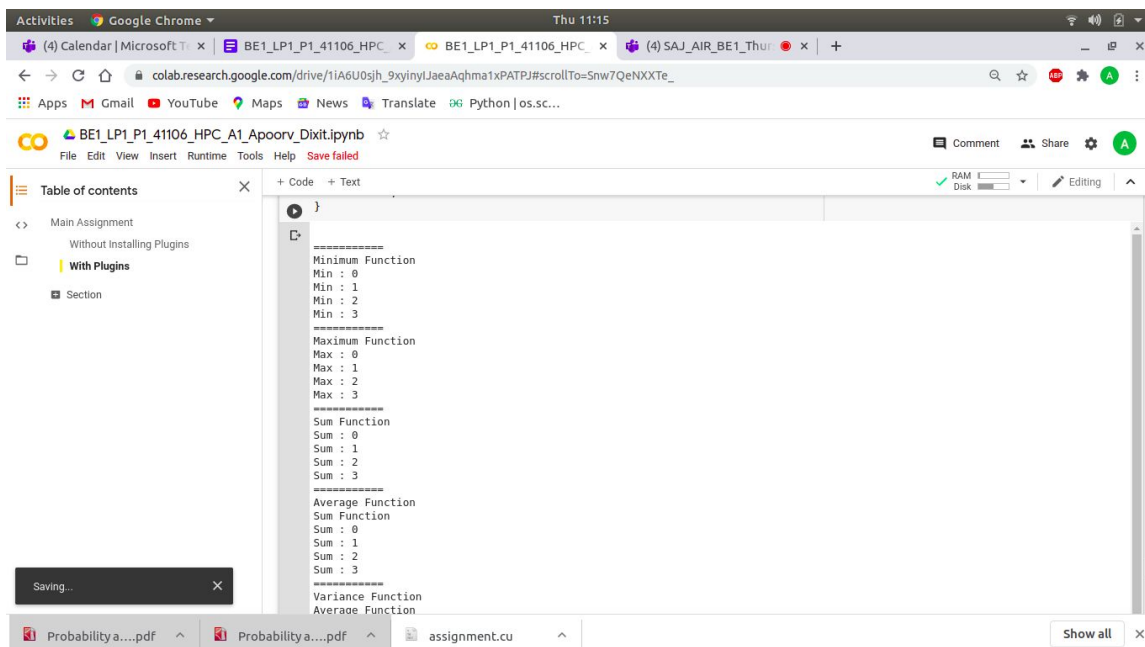
return 0;
}
```

Vector - [383, 386, 277, 415, 293, 335, 386, 492]

Minimum - 277
Maximum - 492
Sum - 2967
Average - 370.875
Variance - 4173.36
Standard Deviation - 64.6015

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

ThreadID and Function Output (Modification)



The screenshot shows a Jupyter Notebook interface with a C++ code cell. The code calculates the minimum, maximum, sum, average, variance, and standard deviation of an array. The output is displayed as a text box below the code cell.

```
Minimum Function
Min : 0
Min : 1
Min : 2
Min : 3

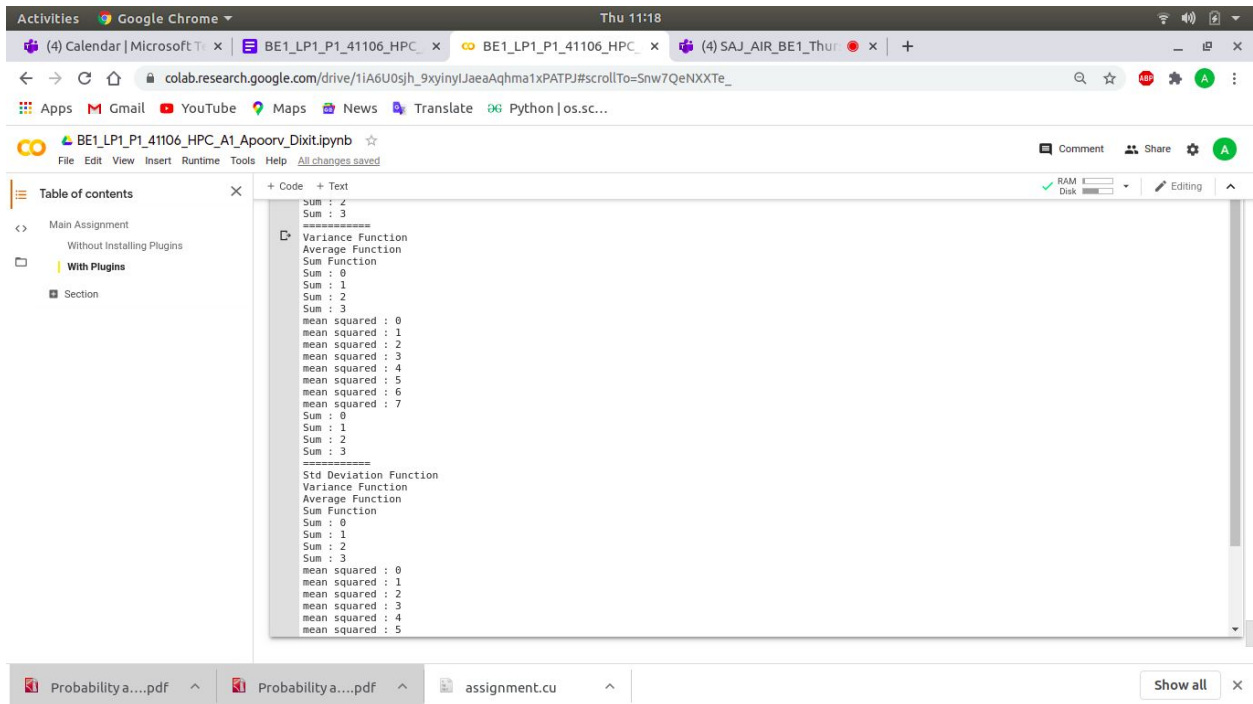
Maximum Function
Max : 0
Max : 1
Max : 2
Max : 3

Sum Function
Sum : 0
Sum : 1
Sum : 2
Sum : 3

Average Function
Sum Function
Sum : 0
Sum : 1
Sum : 2
Sum : 3

Variance Function
Average Function
```

Saving...



Conclusion

I have successfully divided problems into subproblems, learnt how to use GPUs and learnt how to solve sub problems using threads on GPU cores.